

# Auto MPG Prediction

## 1. Project Summary

The goal of this project is to build and compare multiple regression models to predict miles per gallon (MPG) from the Auto MPG dataset, track all experiments with **MLflow**, select the best-performing model, register it in the MLflow Model Registry, and deploy it as a REST API using **Flask**.

We tested **5 regression models**:

1. Linear Regression
2. Ridge Regression
3. Random Forest Regressor
4. Gradient Boosting Regressor
5. XGBoost Regressor

The best model was **Random Forest (version 3 in the MLFlow Registry)**, which was deployed to production.

An example API call returned a prediction of **34.019999 MPG**.

## Step 1 — Data Loading & Preprocessing

**Goal:** Load the dataset, clean it, and prepare training/test splits.

Key tasks:

1. Load `auto-mpg.csv`, treating `?` as missing values.
2. Check and impute missing **horsepower** values with the mean.
3. Drop rows missing the target **mpg**.
4. Separate features (X) and target (y).
5. Identify numeric and categorical columns.
6. Split into training and test sets.

In [2]: df

```
Out[2]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model_year	origin	car_name
0	18.0	8	307.0	130.0	3504.0	12.0	70	1	chevrolet chevelle malibu
1	15.0	8	350.0	165.0	3693.0	11.5	70	1	buick skylark 320
2	18.0	8	318.0	150.0	3436.0	11.0	70	1	plymouth satellite
3	16.0	8	304.0	150.0	3433.0	12.0	70	1	amc rebel sst
4	17.0	8	302.0	140.0	3449.0	10.5	70	1	ford torino
...	...	...	...	...	...	...	...	...	...
393	27.0	4	140.0	86.0	2790.0	15.6	82	1	ford mustang gl
394	44.0	4	97.0	52.0	2130.0	24.6	82	2	vw pickup
395	32.0	4	135.0	84.0	2295.0	11.6	82	1	dodge rampage
396	28.0	4	120.0	79.0	2625.0	18.6	82	1	ford ranger
397	31.0	4	119.0	82.0	2720.0	19.4	82	1	chevy s-10

In [6]: df.isnull().sum()

```
Out[6]: mpg      0
cylinders  0
displacement  0
horsepower  6
weight      0
acceleration  0
model_year  0
origin      0
car_name    0
dtype: int64
```

In [7]: df['horsepower'].fillna(df['horsepower'].mean(), inplace=True)

In [8]: df.isnull().sum()

```
Out[8]: mpg      0
cylinders  0
displacement  0
horsepower  0
weight      0
acceleration  0
model_year  0
origin      0
car_name    0
dtype: int64
```

In [9]: df.columns

```
Out[9]: Index(['mpg', 'cylinders', 'displacement', 'horsepower', 'weight',
              'acceleration', 'model_year', 'origin', 'car_name'],
              dtype=object)
```

```
In [10]: import numpy as np
```

```
In [11]: # Drop rows with missing target
df = df.dropna(subset=['mpg'])

# Separate features and target
X = df.drop(columns=['mpg', 'car_name'])
y = df['mpg']

# Identify numerical and categorical columns
num_cols = X.select_dtypes(include=np.number).columns.tolist()
cat_cols = X.select_dtypes(exclude=np.number).columns.tolist()

# Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
In [12]: print(X_train.shape, X_test.shape)
```

```
(318, 7) (80, 7)
```

## Step 2 — Model Training & MLflow Tracking


**Goal:** Train 5 regression models, log metrics and artifacts with MLflow.

### Models:

- Linear Regression
- Ridge Regression
- Random Forest
- Gradient Boosting
- XGBoost

**To view MLflow UI:**

Visit <http://4.227.222.56:5000>



```

with mlflow.start_run(run_name="LinearRegression"):
    model = LinearRegression()
    model.fit(X_train, y_train)
    predictions = model.predict(X_test)

    rmse = np.sqrt(mean_squared_error(y_test, predictions))
    mae = mean_absolute_error(y_test, predictions)

    mlflow.log_metric("rmse", rmse)
    mlflow.log_metric("mae", mae)
    mlflow.log_param("model_type", "Linear Regression")

    signature = infer_signature(X_train, model.predict(X_train))
    mlflow.sklearn.log_model(
        sk_model=model,
        artifact_path="model",
        signature=signature,
        input_example=X_train,
        registered_model_name="MpgEstimator"
    )

```

2025/08/11 10:29:30 INFO mlflow.tracking.fluent: Experiment with name 'MpgEstimationExperiment' does not exist. Creating a new experiment.

/home/azureuser/anaconda3/lib/python3.9/site-packages/mlflow/types/utils.py:452: UserWarning: Hint: Inferred schema contains integer column(s). Integer columns in Python cannot represent missing values. If your input data contains missing values at inference time, it will be encoded as floats and will cause a schema enforcement error. The best way to avoid this problem is to infer the model schema based on a realistic data sample (training dataset) that includes missing values. Alternatively, you can declare integer columns as doubles (float64) whenever these columns may have missing values. See 'Handling Integers with Missing Values' <<https://www.mlflow.org/docs/latest/models.html#handling-integers-with-missing-values>> for more details.

warnings.warn(

2025/08/11 10:29:30 WARNING mlflow.models.model: `artifact\_path` is deprecated. Please use `name` instead.

Successfully registered model 'MpgEstimator'.

2025/08/11 10:29:34 INFO mlflow.store.model\_registry.abstract\_store: Waiting up to 300 seconds for model version to finish creation. Model name: MpgEstimator, version 1

✎ View run LinearRegression at: <http://4.227.222.56:5000/#/experiments/2/runs/936a27e8d47f430da4ba09d53bca7d6f>

🌱 View experiment at: <http://4.227.222.56:5000/#/experiments/2>

Created version '1' of model 'MpgEstimator'.

Experiments

Search experiments

- ☐ Default
- ☐ FareEstimationExperiment
- ☒ MpgEstimationExperiment

MpgEstimationExperiment

Provide Feedback | Add Description

Share

Runs Models Experimental Evaluation Traces

metrics.rmse < 1 and params.model = "tree"

Time created State: Active Datasets Sort: Created + New run

Columns Group by

Run Name	Created	Dataset	Duration	Source	Models
XGBoost	7 hours ago	-	3.2s	ipykeme...	model
GradientBoosting	7 hours ago	-	2.8s	ipykeme...	model
RandomForest	7 hours ago	-	3.1s	ipykeme...	model
RidgeRegression	7 hours ago	-	2.7s	ipykeme...	model
LinearRegression	7 hours ago	-	4.0s	ipykeme...	model

Show more columns (8 total)

Schema Outputs are identical

Metrics

Show diff only


mae	2.253	2.254	1.593	1.702	1.909
rmse	2.863	2.864	2.144	2.251	2.592

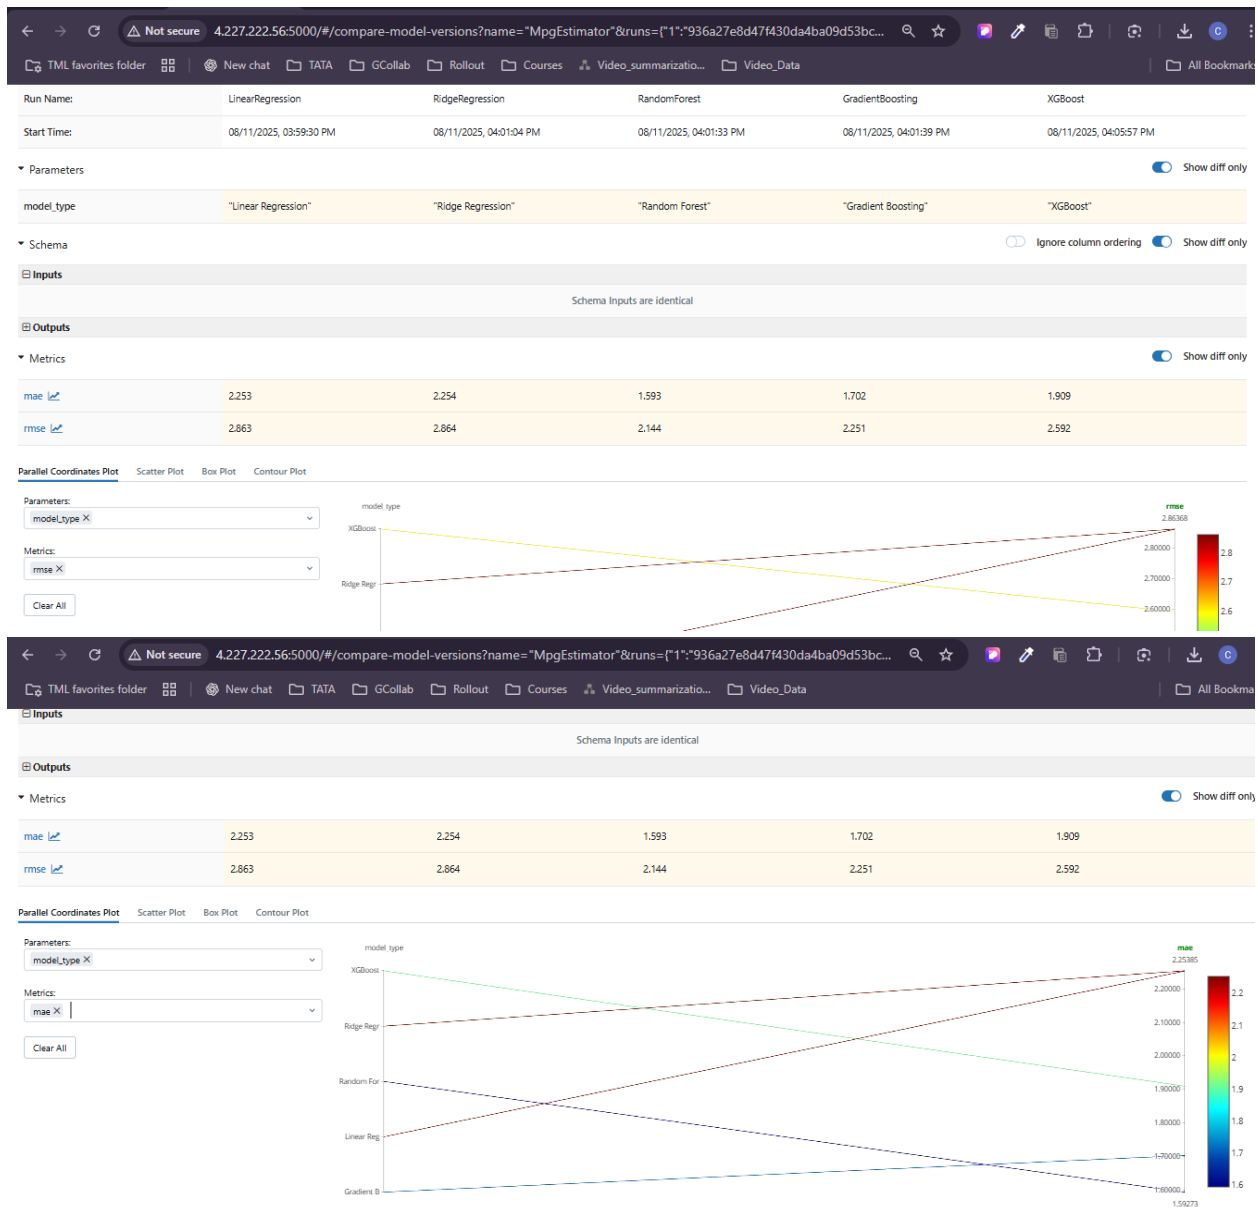
Parallel Coordinates Plot

Parameters: modelType X

Metrics: rmse X

Clear All






## Step 3 — Register the Best Model

After comparing RMSE/MAE in MLflow, Random Forest was chosen the best model as the rmse was the least.

### UI Steps:

1. Open Random Forest run in MLflow.
2. Click **Register Model**, name it RandomForest.
3. Confirm version number (v3).






MpgEstimationExperiment > Models >

 **model**

Overview Traces Artifacts

No description

#### Details

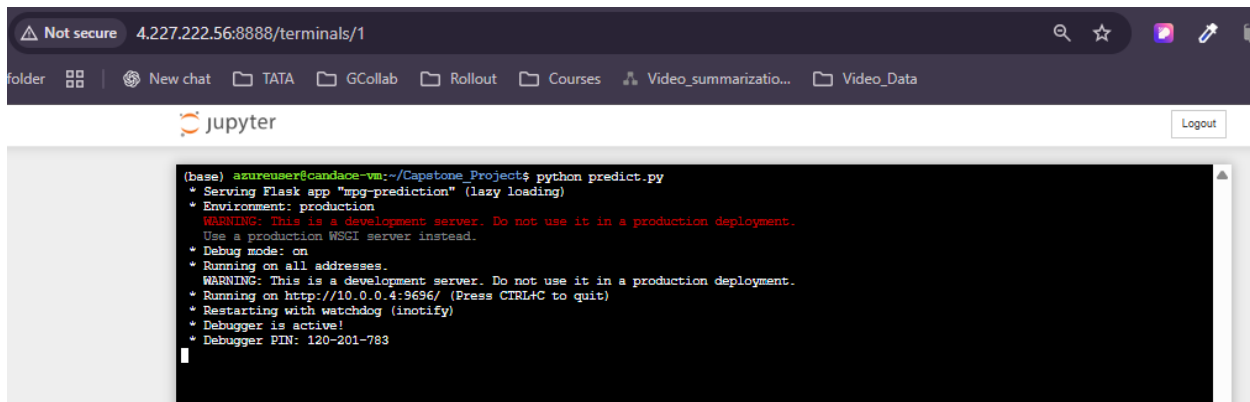
Created at	7 hours ago
Status	 Ready
Model ID	m-c7a7b6de34994ee892dec811b2c38045 
Source run	RandomForest
Source run ID	a0d91a3154d1477f97ce16cc9991d858 
Logged from	 ipykernel_launcher.py
Datasets used	-
Model versions	 MpgEstimator v3 +1

#### Metrics (2)

<input type="text" value="Search metrics"/>			
Metric	Dataset	Source run	Value
rmse	-	RandomForest	2.143801325916186
mae	-	RandomForest	1.5927250000000002

## Step 4 — Deployment with Flask

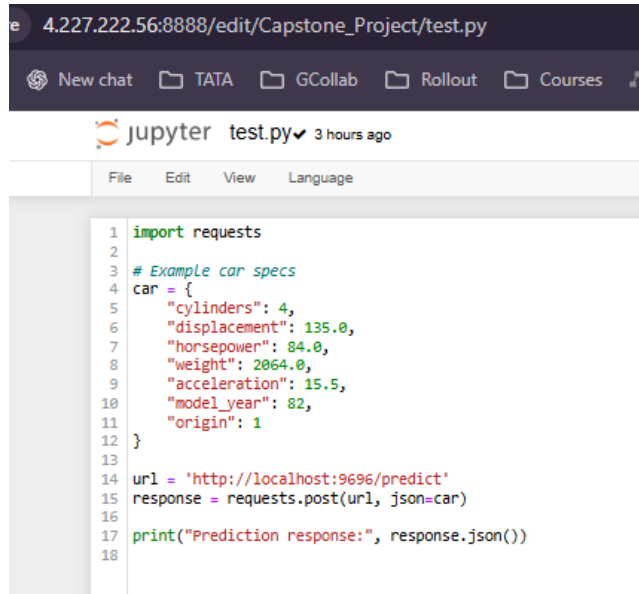
**Goal:** Load the registered model and expose /predict endpoint.



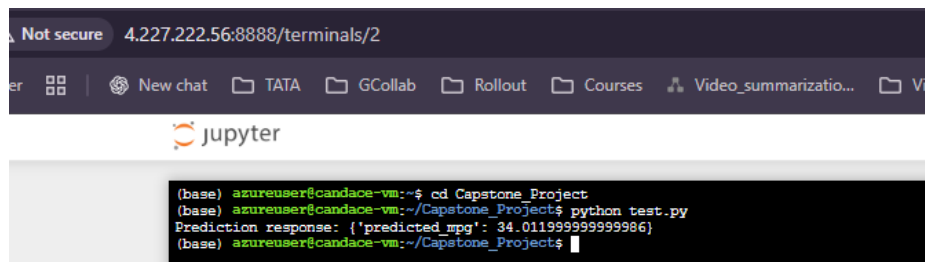
```
(base) azureuser@candace-vm:~/Capstone_Project$ python predict.py
* Serving Flask app "mpg-prediction" (lazy loading)
* Environment: production
  WARNING: This is a development server. Do not use it in a production deployment.
  Use a production WSGI server instead.
* Debug mode: on
* Running on all addresses.
  WARNING: This is a development server. Do not use it in a production deployment.
* Running on http://10.0.0.4:9636/ (Press CTRL+C to quit)
* Restarting with watchdog (inotify)
* Debugger is active!
* Debugger PIN: 120-201-783
```

## Step 5 — API Testing

test.py:



```
1 import requests
2
3 # Example car specs
4 car = {
5     "cylinders": 4,
6     "displacement": 135.0,
7     "horsepower": 84.0,
8     "weight": 2064.0,
9     "acceleration": 15.5,
10    "model_year": 82,
11    "origin": 1
12 }
13
14 url = 'http://localhost:9696/predict'
15 response = requests.post(url, json=car)
16
17 print("Prediction response:", response.json())
18
```



```
(base) azureuser@candace-vm:~$ cd Capstone_Project
(base) azureuser@candace-vm:~/Capstone_Project$ python test.py
Prediction response: {'predicted_mpg': 34.011999999999986}
(base) azureuser@candace-vm:~/Capstone_Project$
```

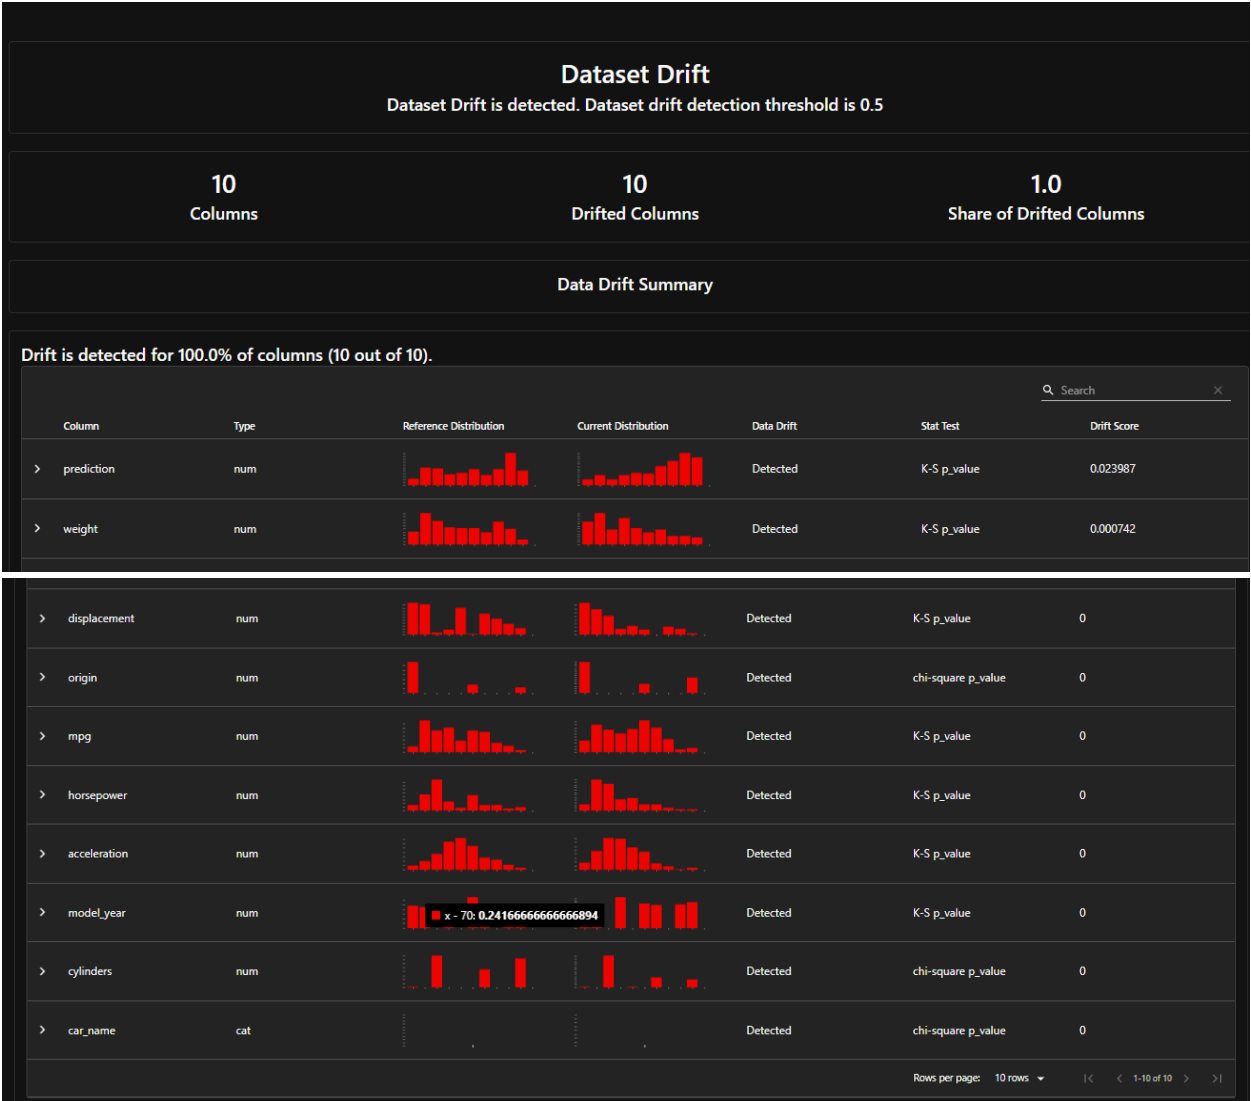
## Results Interpretation

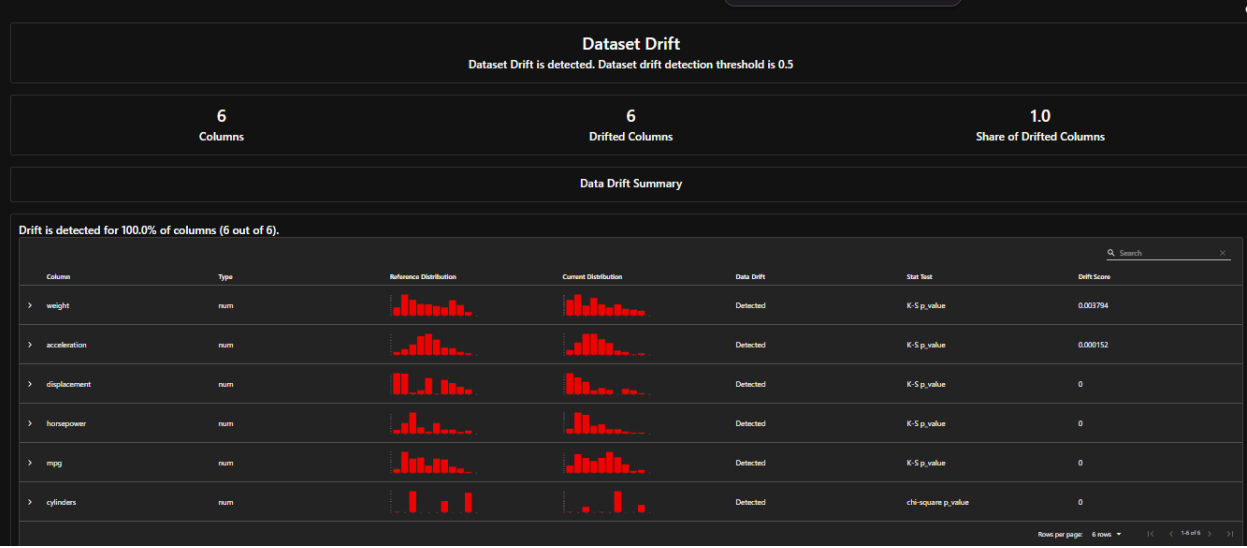
After cleaning the Auto MPG dataset and imputing missing horsepower values, five regression models were trained and evaluated. Metrics were logged with MLflow for reproducibility and comparison. Random Forest achieved the best performance and was registered as version 3 in MLflow. The model was successfully deployed via Flask as a REST API, which returned an MPG prediction of 31.863 for the example input.

## Evidently Data Drift and Model Performance Report:



This script loads and cleans the Auto MPG dataset, splits it into a baseline (reference) and new (current) dataset, and then deliberately alters the current data to simulate drift in key features like weight, horsepower, and acceleration. It trains a simple linear regression model to predict MPG, generates predictions for both datasets, and uses the Evidently library to create HTML reports on data drift and target drift. Finally, it checks the overall drift score from the report against a set threshold and prints an alert if significant drift is detected.





Evidently Model Performance Report

